

# Database Management Systems

Lecture 2

Transactions. Concurrency Control

## Scheduling Transactions

- *schedule*

- a list of operations (Read / Write / Commit / Abort) of  $n$  transactions with the property that the order of the operations in each individual transaction is preserved

# Scheduling Transactions

T1

read(V)  
read(sum)

read(V)  
sum := sum + V  
write(sum)  
commit

T2

read(V)  
V := V + 50  
write(V)  
commit

Schedule

read1(V)  
read1(sum)  
read2(V)

write2(V)  
commit2  
read1(V)

write1(sum)  
commit1

## Serial and Non-Serial Schedules

- serial schedule

- a schedule in which the actions of different transactions are not interleaved

T1

```
read(V)
read(sum)
read(V)
sum := sum + V
write(sum)
commit
```

T2

```
read(V)
V := V + 50
write(V)
commit
```

## Serial and Non-Serial Schedules

- *non-serial schedule*

- a schedule in which the actions of different transactions are interleaved

read1(V)

read1(sum)

read2(V)

write2(V)

commit2

read1(V)

write1(sum)

commit1

## Serializability

- $C$  – set of transactions
- $Sch(C)$  – the set of schedules for  $C$
- serializable schedule
  - a schedule  $S \in Sch(C)$  is serializable if the effect of  $S$  on a consistent database instance is identical to the effect of some serial schedule  $S_0 \in Sch(C)$

# Serializability

- serializability
  - correctness criterion for an interleaved execution schedule
    - consider the serial execution schedule  $(T_1, T_2, \dots, T_n)$ ,  $T_i \in C$ 
      - assume the database instance is in a correct state prior to executing  $T_1$
      - every transaction must preserve the consistency of the database
  - => the database is in a correct state after  $T_n$  completes execution
  - => if a serializable schedule is executed on a correct database instance, it produces a correct database instance (since it is equivalent to some serial schedule)
  - ensuring serializability prevents inconsistencies generated by concurrent transactions that interfere with one another

## Serializability

- serializability
  - objective
    - finding interleaved schedules enabling transactions to execute concurrently without interfering with each other (such schedules result in a correct database state)



## Serializability

- the order of *read* and *write* operations is important
- actions that cannot be swapped in a schedule:
  - actions belonging to the same transaction
  - actions in different transactions operating on the same object if at least one of them is a *write*

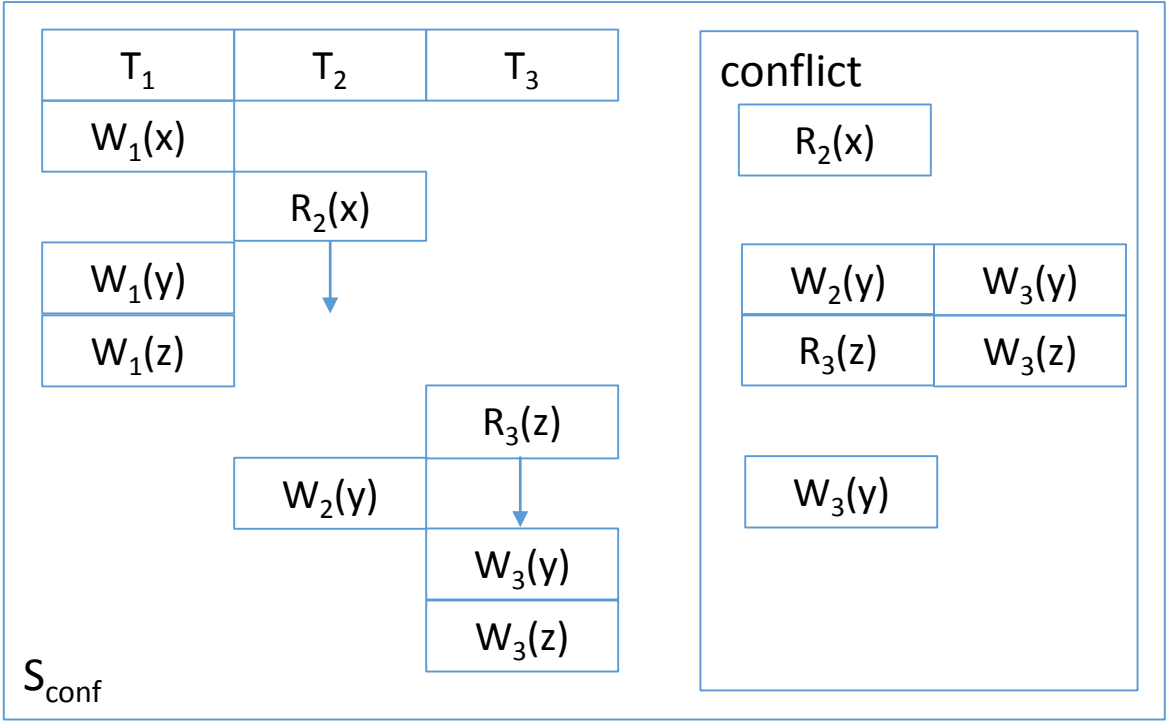
## Conflict Serializability

- $C$  – set of transactions
- $Sch(C)$  – the set of schedules for  $C$
- $Op(C)$  – set of operations of the transactions in  $C$
- consider schedule  $S \in Sch(C)$
- the *conflict relation* of  $S$  is defined as:
  - $conflict(S) = \{(op_1, op_2) \mid op_1, op_2 \in Op(C), op_1 \text{ occurs before } op_2 \text{ in } S, op_1 \text{ and } op_2 \text{ are in conflict}\}$
- two schedules  $S_1$  and  $S_2 \in Sch(C)$  are conflict equivalent, written  $S_1 \equiv_c S_2$ , if  $conflict(S_1) = conflict(S_2)$ , i.e.:
  - $S_1$  and  $S_2$  contain the same operations of the same transactions and
  - every pair of conflicting operations is ordered in the same manner in  $S_1$  and  $S_2$

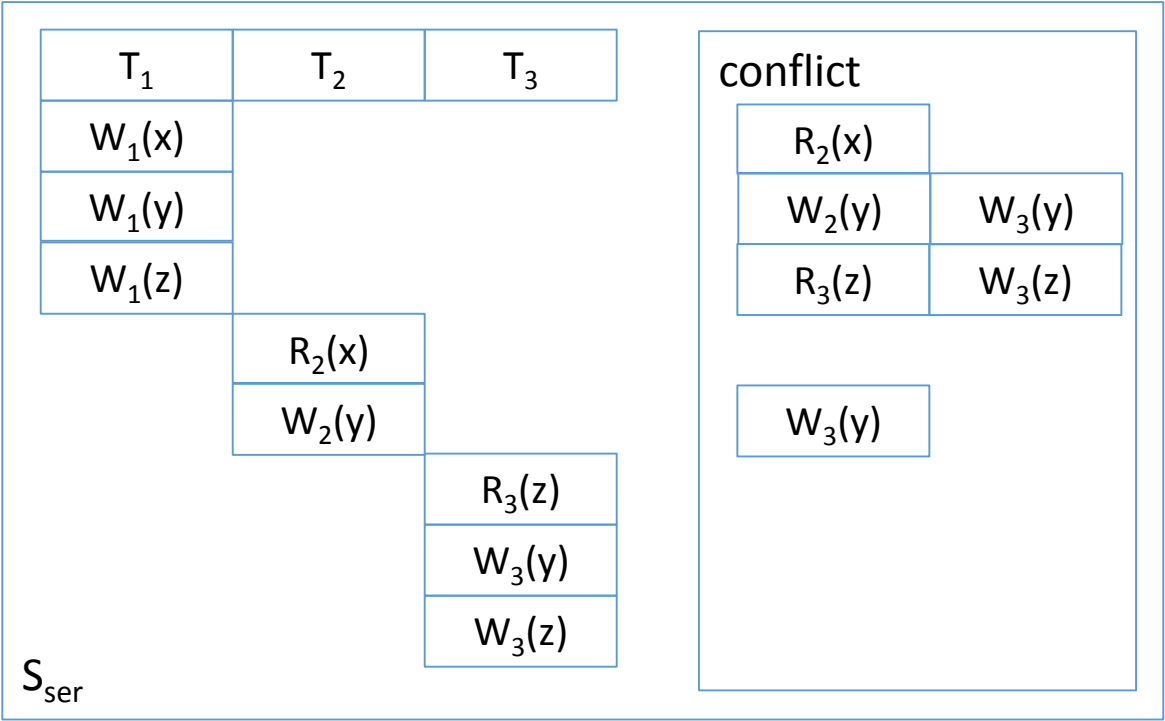
## Conflict Serializability

- $C$  – set of transactions
- $Sch(C)$  – the set of schedules for  $C$
- let  $S$  be a schedule in  $Sch(C)$
- schedule  $S$  is conflict serializable if there exists a serial schedule  $S_0 \in Sch(C)$  such that  $S \equiv_c S_0$ , i.e.,  $S$  is conflict equivalent to some serial schedule

# Conflict Serializability



conflict serializable schedule




serial schedule

## Conflict Serializability - Precedence Graph


- let  $S$  be a schedule in  $Sch(C)$
- the precedence graph (serializability graph) of  $S$  contains:
  - one node for every committed transaction in  $S$
  - an arc from  $T_i$  to  $T_j$  if an action in  $T_i$  precedes and conflicts with one of the actions in  $T_j$
- Theorem:
  - a schedule  $S \in Sch(C)$  is conflict serializable if and only if its precedence graph is acyclic

## Conflict Serializability - Precedence Graph

- example - a schedule that is not conflict serializable:



- the precedence graph has a cycle:




## Conflict Serializability - Precedence Graph


- algorithm to test the conflict serializability of a schedule  $S \in Sch(C)$ 
  1. create a node labeled  $T_i$  in the precedence graph for every committed transaction  $T_i$  in the schedule
  2. create an arc  $(T_i, T_j)$  in the precedence graph if  $T_j$  executes a Read(A) after a Write(A) executed by  $T_i$
  3. create an arc  $(T_i, T_j)$  in the precedence graph if  $T_j$  executes a Write(A) after a Read(A) executed by  $T_i$
  4. create an arc  $(T_i, T_j)$  in the precedence graph if  $T_j$  executes a Write(A) after a Write(A) executed by  $T_i$
  5.  $S$  is conflict serializable if and only if the resulting precedence graph has no cycles

## Conflict Serializability - Precedence Graph

- examples
- let  $S_1$  be a schedule over  $\{T_1, T_2, T_3\}$ :




- the precedence graph for  $S_1$ :






## Conflict Serializability - Precedence Graph

- examples
- let  $S_2$  be a schedule over  $\{T_1, T_2, T_3\}$ :

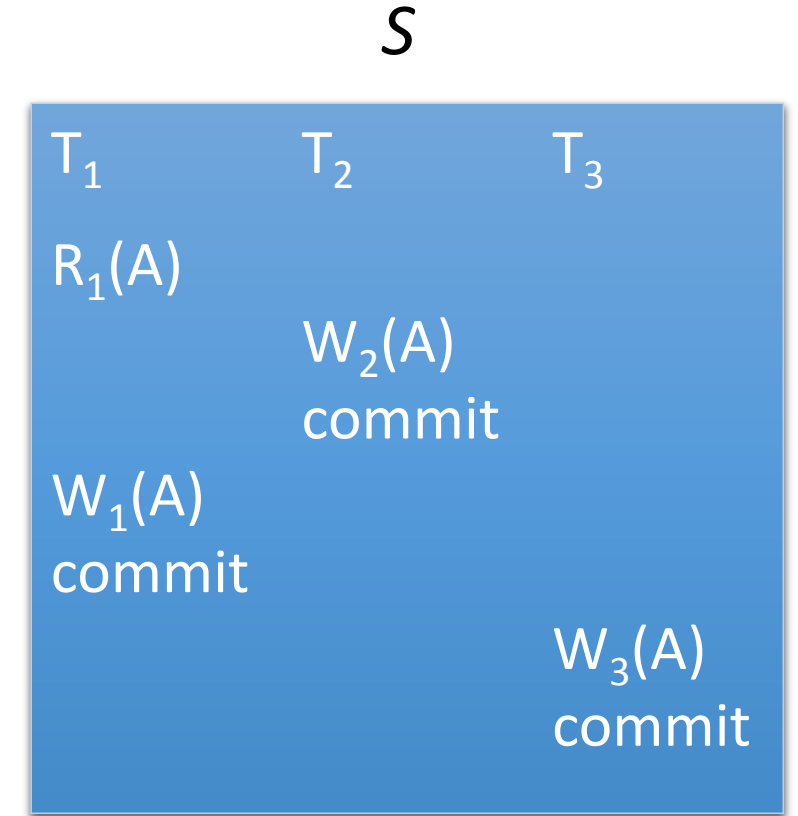


- the precedence graph for  $S_2$ :



## Conflict Serializability

- every conflict serializable schedule is serializable (in the absence of inserts / deletes, when items can only be updated)
- there are serializable schedules that are not conflict serializable
- $S$  is equivalent to the serial execution of transactions  $T_1, T_2, T_3$  (in this order), but it is not conflict equivalent to this serial schedule (the write operations in  $T_1$  and  $T_2$  are ordered differently)



## View Serializability

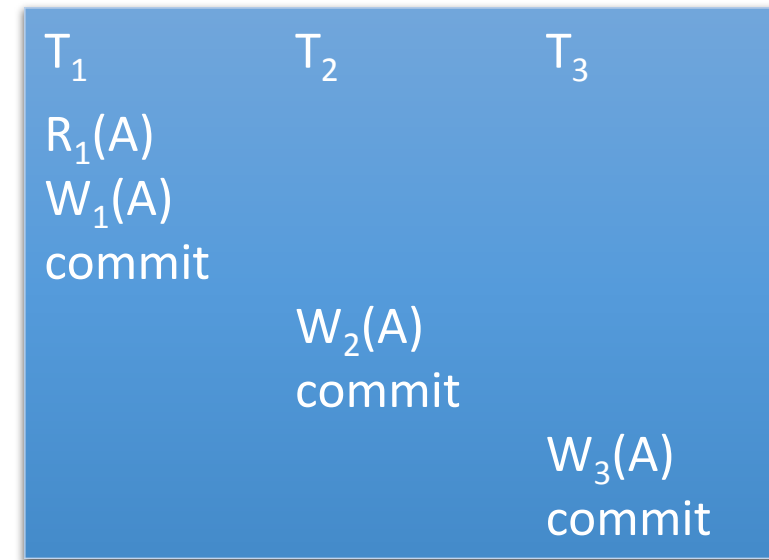
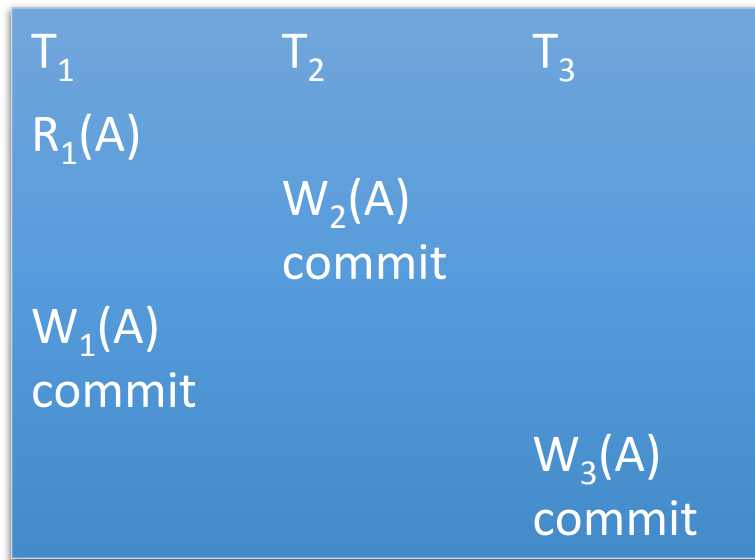
- conflict serializability is a sufficient condition for serializability, but it is not a necessary one
- *view serializability*
  - a more general, sufficient condition for serializability
  - based on *view-equivalence*, a less stringent form of equivalence

## View Serializability

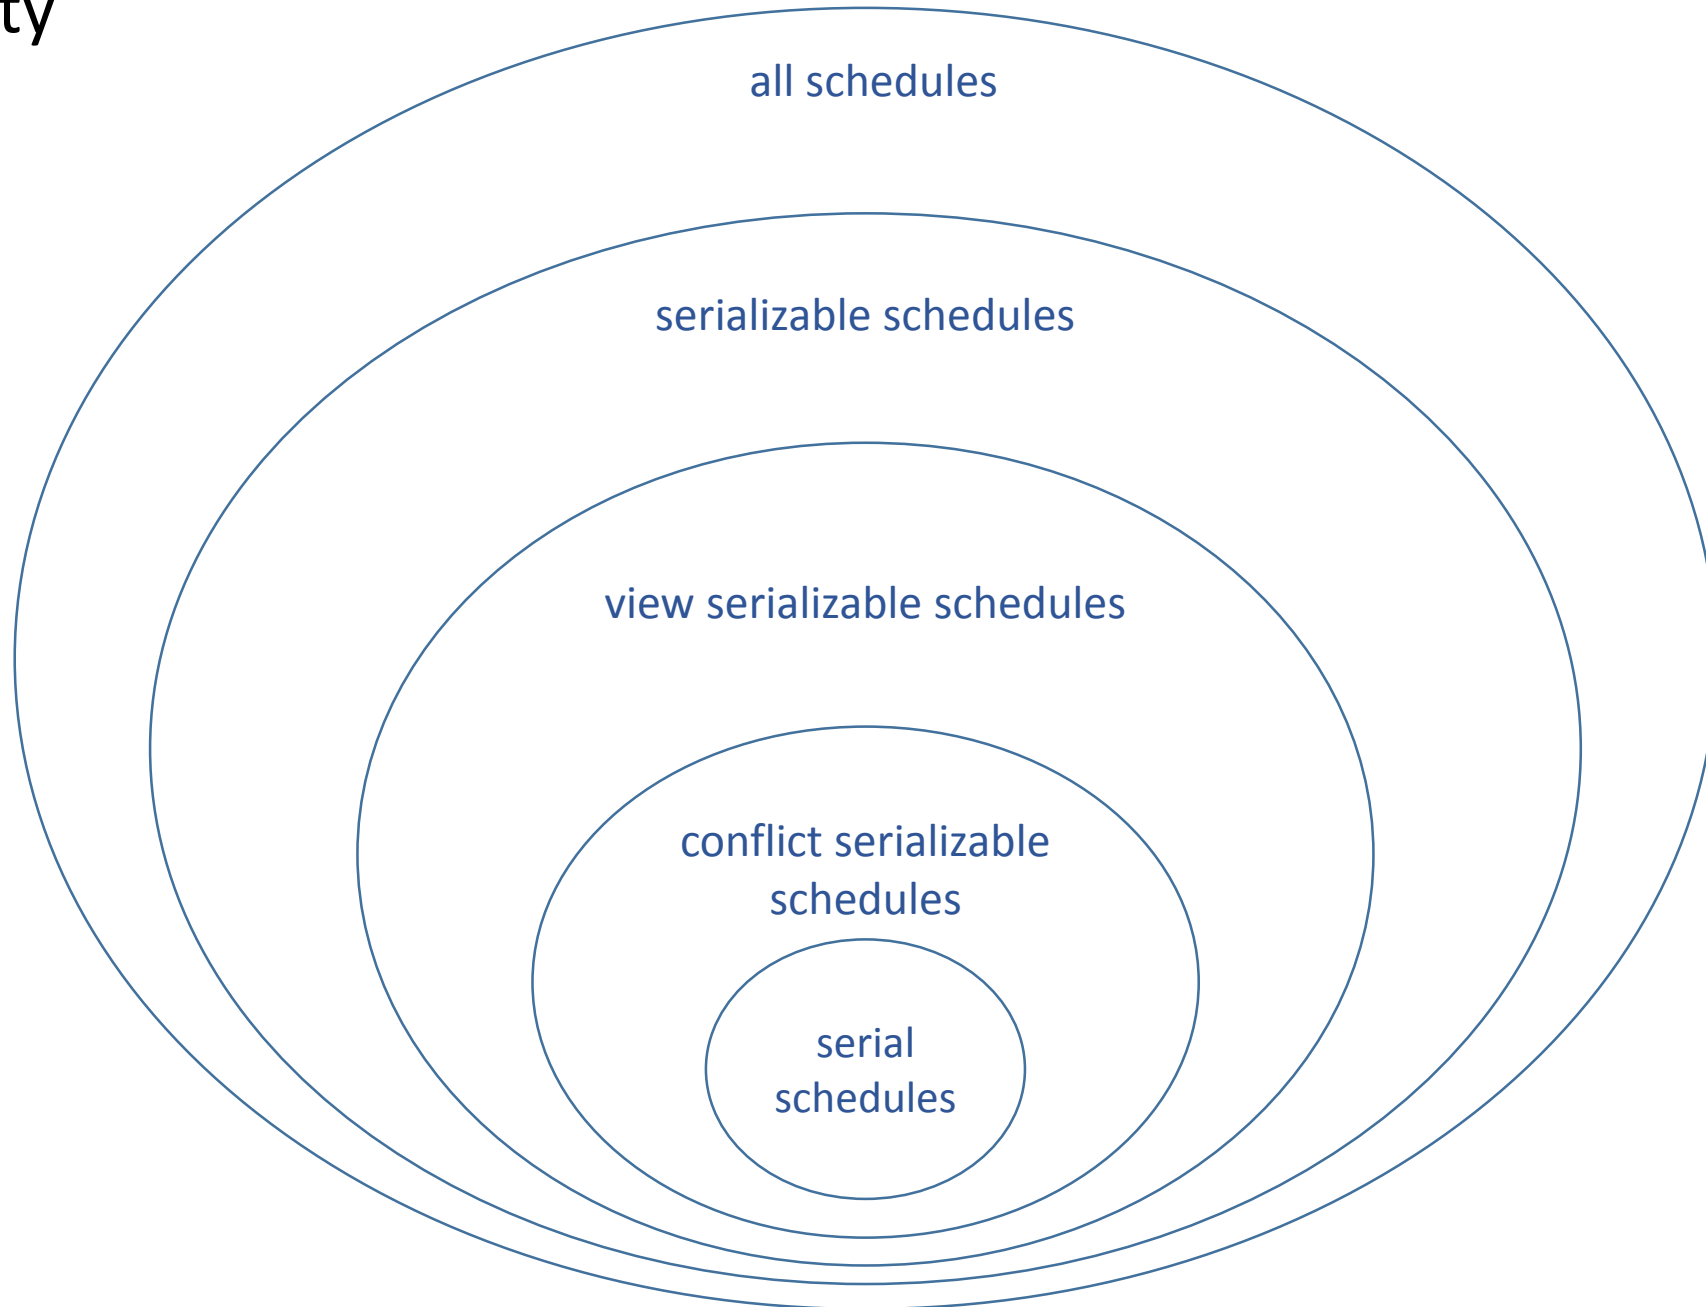
- $C$  – set of transactions
- $Sch(C)$  – the set of schedules for  $C$
- let  $T_i, T_j \in C, S_1, S_2 \in Sch(C)$ ;  $S_1$  and  $S_2$  are view equivalent, written  $S_1 \equiv_v S_2$ , if the following conditions are met:
  - if  $T_i$  reads the initial value of  $V$  in  $S_1$ , then  $T_i$  also reads the initial value of  $V$  in  $S_2$ ;
  - if  $T_i$  reads the value of  $V$  written by  $T_j$  in  $S_1$ , then  $T_i$  also reads the value of  $V$  written by  $T_j$  in  $S_2$ ;
  - if  $T_i$  writes the final value of  $V$  in  $S_1$ , then  $T_i$  also writes the final value of  $V$  in  $S_2$ .
- i.e.:
  - each transaction performs the same computation in  $S_1$  and  $S_2$
  - and
  - $S_1$  and  $S_2$  produce the same final database state.

## View Serializability

- $C$  – set of transactions
- $Sch(C)$  – the set of schedules for  $C$
- a schedule  $S \in Sch(C)$  is view serializable if there exists a serial schedule  $S_0 \in Sch(C)$  such that  $S \equiv_v S_0$ , i.e.,  $S$  is view equivalent to some serial schedule

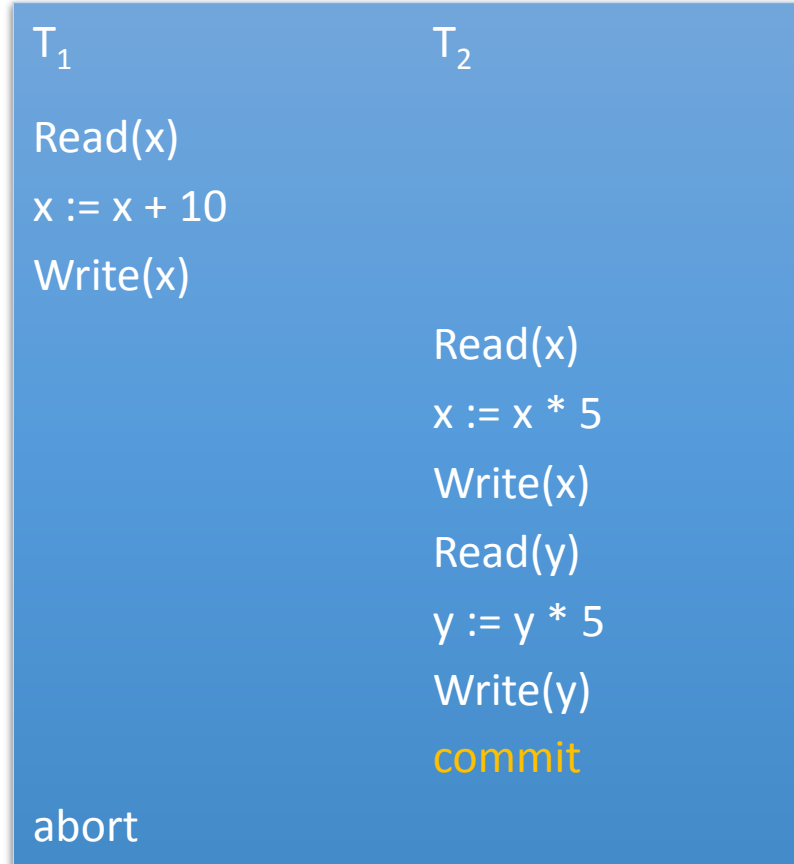


# Serializability



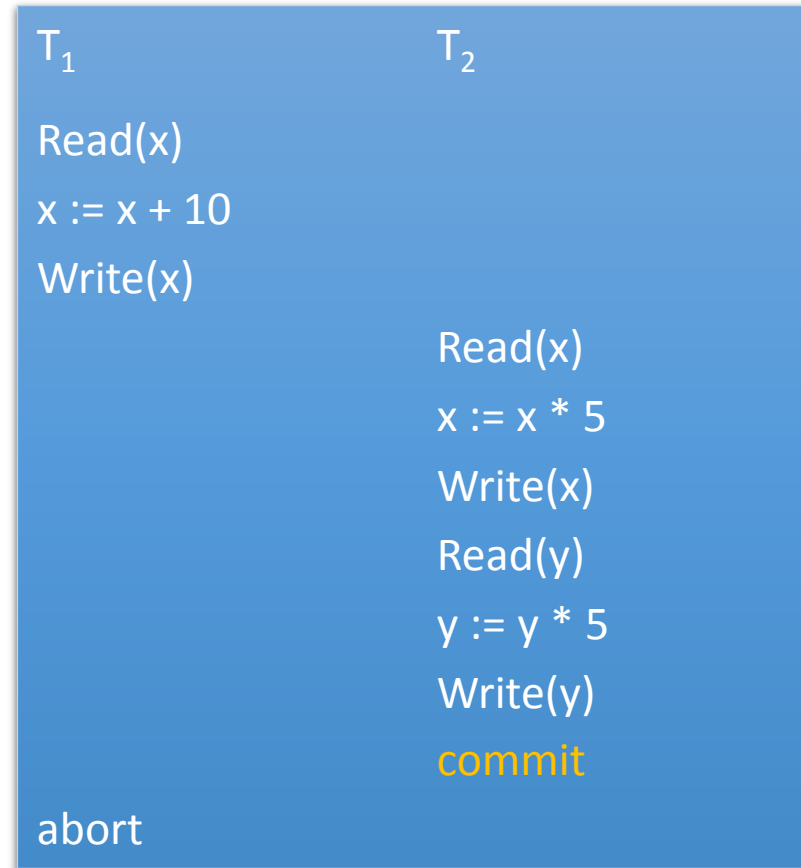
## Recoverable Schedules

- consider schedule  $S$  over  $\{T_1, T_2\}$



- $T_2$  operates on a value of  $x$  that shouldn't have been in the database, since  $T_1$  aborted

# Recoverable Schedules



- cannot cascade the abort of  $T_1$ , since  $T_2$  has already committed
- schedule  $S$  is *unrecoverable*

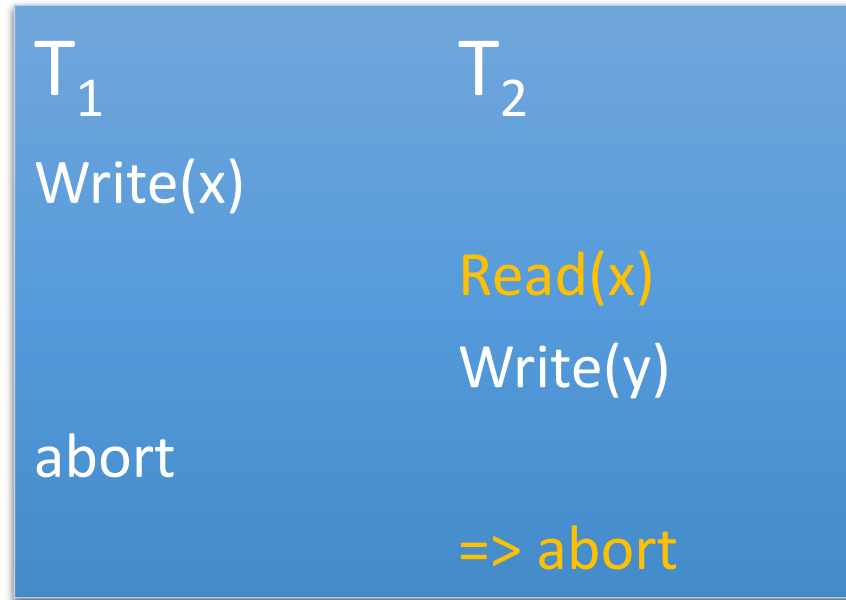


## Recoverable Schedules

- recoverable schedule

- a schedule in which a transaction  $T$  commits only after all transactions whose changes  $T$  read commit

## Avoiding Cascading Aborts



- a schedule in which a transaction  $T$  is reading only changes of committed transactions is said to avoid cascading aborts
- avoiding cascading aborts => recoverable schedules

## Lock-Based Concurrency Control

- technique used to guarantee serializable, recoverable schedules
- *lock*
  - a tool used by the transaction manager to control concurrent access to data
  - prevents a transaction from accessing a data object while another transaction is accessing the object
- *transaction protocol*
  - a set of rules enforced by the transaction manager and obeyed by all transactions
  - example – simple protocol: before a transaction can read / write an object, it must acquire an appropriate lock on the object
  - locks in conjunction with transaction protocols allow interleaved executions

## Lock-Based Concurrency Control

- *transaction protocol*
  - it's impractical for the DBMS to test the serializability of schedules, since the operating system could determine the interleaving of operations
  - instead, the DBMS uses protocols known to produce serializable schedules

## Lock-Based Concurrency Control

- SLock (*shared or read lock*)
  - if a transaction holds an SLock on an object, it can read the object, but it cannot modify it
- XLock (*exclusive or write lock*)
  - if a transaction holds an XLock on an object, it can both read and write the object
- if a transaction holds an SLock on an object, other transactions can be granted SLocks on the object, but they cannot acquire XLocks on it
- if a transaction holds an XLock on an object, other transactions cannot be granted either SLocks or XLocks on the object

	Shared	Exclusive
Shared	Yes	No
Exclusive	No	No

## Lock-Based Concurrency Control

- lock upgrade
  - an SLock granted to a transaction can be upgraded to an XLock
- transactions are issuing lock requests to the lock manager
- locks are held until being explicitly released by transactions
- lock acquire / lock release requests are automatically inserted into transactions by the DBMS (not the user's responsibility)
- locking / unlocking
  - atomic operations

## Lock-Based Concurrency Control

- *lock table*

- structure used by the lock manager to keep track of granted locks / lock requests
- entry in the lock table (corresponding to one data object):
  - number of transactions holding a lock on the data object
  - lock type (SLock / XLock)
  - pointer to a queue of lock requests


## Lock-Based Concurrency Control

- *transactions table*
  - structure maintained by the DBMS
  - one entry / transaction
  - keeps a list of locks held by every transaction



## Strict Two-Phase Locking (Strict 2PL)

- all the locks held by a transaction are released when it completes execution
- the Strict 2PL protocol allows only serializable schedules



# Strict Two-Phase Locking

$T_1$   
R(CM)  
W(CM)

R(CS)  
W(CS)  
commit

$T_2$   
R(CM)  
W(CM)  
R(CS)  
W(CS)  
commit

$T_1$   
XLock(CM)  
R(CM)  
W(CM)  
XLock(CS)  
R(CS)  
W(CS)  
commit

$T_2$   
  
  
  
  
  
  
XLock(CM)  
R(CM)  
W(CM)  
XLock(CS)  
R(CS)  
W(CS)  
commit

# Strict Two-Phase Locking

$T_1$

XLock(CM)

R(CM)

W(CM)

XLock(CS)

R(CS)

W(CS)

commit

$T_2$

XLock(CT)

R(CT)

W(CT)


commit

# Deadlocks

- lock-based concurrency control techniques can lead to deadlocks
- deadlock
  - cycle of transactions waiting for one another to release a locked resource
  - normal execution can no longer continue without an external intervention, i.e., deadlocked transactions cannot proceed until the deadlock is resolved
- deadlock management
  - deadlock prevention
  - deadlock detection
    - allow deadlocks to occur and resolve them when they arise

# Deadlocks

- example



```
T1
BEGIN TRAN
XLock(V)
Read(V)
V := V + 100
Write(V)
XLock(Y)
Wait
Wait
...
```

```
T2
BEGIN TRAN
XLock(Y)
Read(Y)
Y = Y * 5
Write(Y)
XLock(V)
Wait
Wait
...
```

## Deadlocks - Prevention

- assign transactions timestamp-based priorities
  - i.e., the older a transaction is, the higher its priority
- 2 policies: Wait-Die and Wound-Wait
- assume  $T_i$  wants to access an object locked by  $T_j$ 
  - Wait-Die
    - if  $T_i$ 's priority is higher,  $T_i$  can wait; otherwise,  $T_i$  is aborted
  - Wound-Wait
    - if  $T_i$ 's priority is higher,  $T_j$  is aborted; otherwise,  $T_i$  can wait
- if an aborted transaction is restarted, it's assigned its original timestamp

## Deadlocks - Prevention

- example – Wait-Die / Wound-Wait

		Wait-die		Wound-wait	
T1	T2	T1	T2	T1	T2
-----		-----		-----	
R1(x)		R1(x)		R1(x)	
	R2(y)		R2(y)		R2(y)
W1(y)		W1(y) (wait)		W1(y)	(Abort T2)
	W2(x)		W2(x) (A)		

# References

- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2<sup>nd</sup> Edition), McGraw-Hill, 2000
- [Le99] LEVENE, M., LOIZOU, G., A Guided Tour of Relational Databases and Beyond, Springer, 1999
- [Da03] DATE, C.J., An Introduction to Database Systems (8<sup>th</sup> Edition), Addison-Wesley, 2003
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007, <http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si10] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, McGraw-Hill, 2010, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems, <http://infolab.stanford.edu/~ullman/fcdb.html>