

# Database Management Systems

Lecture 3

Transactions. Concurrency Control (II)

## Deadlocks - Detection




### a. waits-for graph

- node / active transaction
- arc from  $T_i$  to  $T_j$  if  $T_i$  is waiting for  $T_j$  to release a lock
- cycle in the graph  $\Rightarrow$  deadlock
- DBMS periodically checks whether there are cycles in the waits-for graph

# Deadlocks - Detection

## a. waits-for graph

- example




# Deadlocks - Detection

## a. waits-for graph

- example

| $T_1$ | $T_2$ | $T_3$ | $T_4$ |
|-------|-------|-------|-------|
| S(A)  | S(A)  | S(A)  |       |
| R(A)  | R(A)  | R(A)  |       |
|       |       |       | S(B)  |
|       |       |       | R(B)  |
| X(B)  | X(B)  | X(B)  |       |
|       |       |       | X(A)  |
| ...   | ...   | ...   | ...   |



## Deadlocks - Detection

### b. timeout mechanism


- very simple, practical method of detecting deadlocks used by several DBMSs
- if a transaction T has been waiting too long for a lock on an object, a deadlock is assumed to exist and T is terminated

## Deadlocks – Choosing the Deadlock Victim

- possible criteria to consider when choosing the deadlock victim
  - the number of objects modified by the transaction
  - the number of objects that are to be modified by the transaction
  - the number of locks held
- the policy should be “fair”, i.e., if a transaction is repeatedly chosen as a victim, it should be eventually allowed to proceed

## Two-Phase Locking (2PL)

- once a transaction releases a lock, it cannot request other locks
- phase 1
  - *growing phase*
  - transaction acquires locks
- phase 2
  - *shrinking phase*
  - transaction releases locks



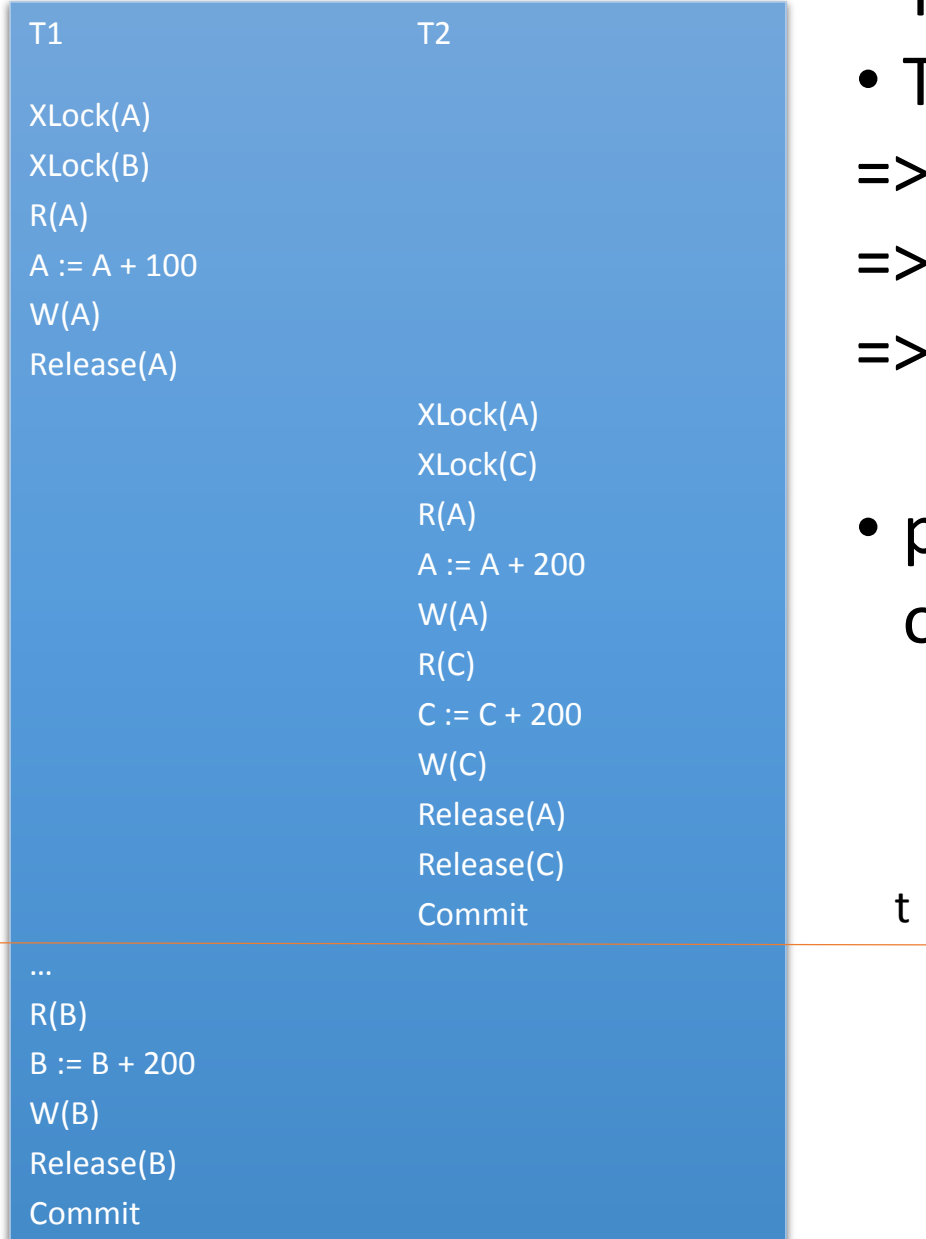
## Two-Phase Locking

- $C$  – set of transactions
- $Sch(C)$  – set of schedules for  $C$
- if all transactions in  $C$  obey 2PL, then any schedule  $S \in Sch(C)$  that completes normally is serializable



# Two-Phase Locking

- example



- T1, T2 – 2PL
- T1 forced to terminate at time t
  - => T1's updates are undone
  - => T2's update to A is lost
  - => atomicity is compromised
- problem – T1 release its exclusive lock on A prior to completion

## Strict Schedules

- transaction  $T_i$  wrote object  $A \Rightarrow$  transaction  $T_j$  can read / write  $A$  only after  $T_i$ 's completion (commit / abort)
- strict schedules  $\Rightarrow$  avoiding cascading aborts  $\Rightarrow$  recoverable schedules
- Strict 2PL only allows strict schedules

# The Phantom Problem

example 1. Researchers[RID, ..., ImpactFactor, Age]

- Page1: <R1, 5, 30>, <R2, 5, 20>
- Page2: <R4, 5, 100>
- Page3: <R8, 6, 18>, <R9, 6, 19>
- concurrent transactions T1 and T2
- transaction T1
  - retrieve age of oldest researcher for impact factor values 5 and 6
- transaction T2
  - add new researcher with impact factor 5
  - remove researcher R9
- T1 and T2 obey Strict 2PL

## The Phantom Problem

- Page1: <R1, 5, 30>, <R2, 5, 20>
- Page2: <R4, 5, 100>
- Page3: <R8, 6, 18>, <R9, 6, 19>
- outcome of interleaved schedule on the right:
  - IF 5, Max Age 100
  - IF 6, Max Age 18
- outcome of serial schedule (T1T2):
  - IF 5, Max Age 100
  - IF 6, Max Age 19
- outcome of serial schedule (T2T1):
  - IF 5, Max Age 102
  - IF 6, Max Age 18

| T1                               | T2                               |
|----------------------------------|----------------------------------|
| XLock(Page1)                     |                                  |
| XLock(Page2)                     |                                  |
| compute max age for IF 5 => 100  |                                  |
|                                  | XLock(Page4)                     |
|                                  | XLock(Page3)                     |
|                                  | add record <R5, 5, 102> on Page4 |
|                                  | delete researcher R9             |
|                                  | commit – all locks are released  |
| XLock(Page3)                     |                                  |
| compute max age for IF = 6 => 18 |                                  |
| ...                              |                                  |

## The Phantom Problem

=> the interleaved schedule is not serializable, as its outcome is not identical to the outcome of any serial schedule

- however, the schedule is conflict serializable (acyclic precedence graph)

=> in the presence of insert operations, i.e., if new objects can be added to the database, conflict serializability does not guarantee serializability

# The Phantom Problem

example 2.

- T1 executes the same query twice
- between the 2 read operations, another transaction T2 inserts a row that meets the condition in T1's query; T2 commits

T1  
SELECT \*  
FROM Students  
WHERE GPA >= 7

T2

INSERT INTO Students VALUES  
(23, 'Ioana', 'Gheorghiu', 10)  
COMMIT

SELECT \*  
FROM Students  
WHERE GPA >= 7

...

## Isolation Levels \*

- determine the degree to which a transaction is isolated from the changes made by other concurrently running transactions
- greater concurrency -> concurrency anomalies
- 4 isolation levels
  - READ UNCOMMITTED
  - READ COMMITTED
  - REPEATABLE READ
  - SERIALIZABLE

\* check seminar 3

## Isolation Levels

- READ UNCOMMITTED

- a transaction must acquire an exclusive lock prior to writing an object
- no locks are requested when reading objects
- exclusive locks are released at the end of the transaction
- lowest degree of isolation



# Isolation Levels

- READ UNCOMMITTED

- dirty reads ✓

- unrepeatable reads ✓

- phantoms ✓

| T1   | T2   |
|------|------|
| R(A) |      |
| W(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(B) |      |
| W(B) |      |
| A    |      |

| T1   | T2   |
|------|------|
| R(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(A) |      |
| W(A) |      |
| C    |      |

| T1  | T2                            |
|---|-------------------------------|
| R(students with<br>id between 100<br>and 110) |                               |
|   | insert student with<br>id 101 |
|   | C                             |
| R(students with<br>id between 100<br>and 110) |                               |
| C   |                               |

## Isolation Levels

- READ COMMITTED

- a transaction must acquire an exclusive lock prior to writing an object
- a transaction must acquire a shared lock prior to reading an object (i.e., the last transaction that modified the object is finished)
- exclusive locks are released at the end of the transaction
- shared locks are immediately released

# Isolation Levels

- READ COMMITTED

- dirty reads ✗

| T1   | T2   |
|------|------|
| R(A) |      |
| W(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(B) |      |
| W(B) |      |
| A    |      |

- unrepeatable reads ✓

| T1   | T2   |
|------|------|
| R(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(A) |      |
| W(A) |      |
| C    |      |

- phantoms ✓

| T1  | T2                            |
|---|-------------------------------|
| R(students with<br>id between 100<br>and 110) |                               |
|   | insert student with<br>id 101 |
|   | C                             |
| R(students with<br>id between 100<br>and 110) |                               |
| C   |                               |

## Isolation Levels

- REPEATABLE READ

- a transaction must acquire an exclusive lock prior to writing an object
- a transaction must acquire a shared lock prior to reading an object
- exclusive locks are released at the end of the transaction
- shared locks are released at the end of the transaction

# Isolation Levels

- REPEATABLE READ

- dirty reads ✖

| T1   | T2   |
|------|------|
| R(A) |      |
| W(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(B) |      |
| W(B) |      |
| A    |      |

- unrepeatable reads ✖

| T1   | T2   |
|------|------|
| R(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(A) |      |
| W(A) |      |
| C    |      |

- phantoms ✓

| T1  | T2                            |
|---|-------------------------------|
| R(students with<br>id between 100<br>and 110) |                               |
|   | insert student with<br>id 101 |
|   | C                             |
| R(students with<br>id between 100<br>and 110) |                               |
| C   |                               |

## Isolation Levels

- SERIALIZABLE

- a transaction must acquire locks on objects before reading / writing them
- a transaction can also acquire locks on sets of objects that must remain unmodified
- locks are held until the end of the transaction
- highest degree of isolation

# Isolation Levels

- SERIALIZABLE

- dirty reads

✗

unrepeatable reads ✗

phantoms ✗

| T1   | T2   |
|------|------|
| R(A) |      |
| W(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(B) |      |
| W(B) |      |
| A    |      |

| T1   | T2   |
|------|------|
| R(A) |      |
|      | R(A) |
|      | W(A) |
|      | C    |
| R(A) |      |
| W(A) |      |
| C    |      |

| T1  | T2                            |
|---|-------------------------------|
| R(students with<br>id between 100<br>and 110) |                               |
|   | insert student with<br>id 101 |
|   | C                             |
| R(students with<br>id between 100<br>and 110) |                               |
| C   |                               |

# References

- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2<sup>nd</sup> Edition), McGraw-Hill, 2000
- [Le99] LEVENE, M., LOIZOU, G., A Guided Tour of Relational Databases and Beyond, Springer, 1999
- [Da03] DATE, C.J., An Introduction to Database Systems (8<sup>th</sup> Edition), Addison-Wesley, 2003
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,  
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si10] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, McGraw-Hill, 2010, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,  
<http://infolab.stanford.edu/~ullman/fcdb.html>