

# Database Management Systems

Lecture 9

Evaluating Relational Operators

Query Optimization (III)

- running example - schema
  - Students (SID: integer, SName: string, Age: integer)
  - Courses (CID: integer, CName: string, Description: string)
  - Exams (SID: integer, CID: integer, EDate: date, Grade: integer, FacultyMember: string)
- Students
  - every record has 50 bytes
  - there are 80 records / page
  - 500 pages
- Courses
  - every record has 50 bytes
  - there are 80 records / page
  - 100 pages
- Exams
  - every record has 40 bytes
  - there are 100 records / page
  - 1000 pages

# Selection

Q:

```
SELECT *  
FROM Exams E  
WHERE E.FacultyMember = 'Ionescu'
```

- use information in the selection condition to reduce the number of retrieved tuples
- e.g.,  $|Q| = 4$ , B+ tree index on FacultyMember
  - it's expensive to scan E (1000 I/Os) to evaluate the query
  - should use the index instead
- selection algorithms - techniques
  - iteration, indexing

\* E - M pages,  $p_E$  records / page \*      \* 1000 pages \*      \* 100 records / page\*

## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- no index on *attr*, data not sorted on *attr*
  - must scan E and test the condition for each tuple
  - access path: file scan

=> cost: M I/Os = 1000 I/Os
- no index, sorted data (E physically sorted on *attr*)
  - \* binary search to locate 1<sup>st</sup> tuple that satisfies condition and
  - \* scan E starting at this position until condition is no longer satisfied
- access method: sorted file scan

## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- no index, sorted data (E physically sorted on *attr*)  
=> cost:
  - binary search:  $O(\log_2 M)$
  - scan cost: varies from 0 to M
  - binary search on E
    - $\log_2 1000 \approx 10$  I/Os

## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- B+ tree index on *attr*
  - \* search tree to find 1<sup>st</sup> index entry pointing to a qualifying E tuple
    - cost: typically 2, 3 I/Os
  - \* scan leaf pages to retrieve all qualifying entries
    - cost: depends on the number of qualifying entries
  - \* for each qualifying entry - retrieve corresponding tuple in E
    - cost: depends on the number of tuples and the nature of the index (clustered / non-clustered)

## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- B+ tree index on *attr*
  - assumption
    - indexes use a2 or a3
    - a1-based index => data entry contains the data record => the cost of retrieving records = the cost of retrieving the data entries!
  - access path: B+ tree index
    - clustered index:
      - best access path when *op* is not *equality*
      - good access path when *op* is *equality*

## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- B+ tree index on *attr*

Q

```
SELECT *  
FROM Exams E  
WHERE E.FacultyMember < 'C%'
```

- names uniformly distributed with respect to 1<sup>st</sup> letter

=> |Q|  $\approx$  10,000 tuples = 100 pages

- clustered B+ tree index on FacultyMember

=> cost of retrieving tuples:  $\approx$  100 I/Os (a few I/Os to get from root to leaf)

- non-clustered B+ tree index on FacultyMember

=> cost of retrieving tuples: up to 1 I/O per tuple (worst case) => up to 10.000 I/Os



## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$

- B+ tree index on *attr*

SELECT \*

FROM Exams E

WHERE E.FacultyMemger < 'C%'

- refinement - sort rids in qualifying data entries by page-id  
=> a page containing qualifying tuples is retrieved only once
  - cost of retrieving tuples: number of pages containing qualifying tuples (but such tuples are probably stored on more than 100 pages)
- range selections
  - non-clustered indexes can be expensive
  - could be less costly to scan the relation (in our example: 1000 I/Os)

## Selection

- general selections
  - selections without disjunctions
- C - CNF condition without disjunctions
  - evaluation options:
    1. use the most selective access path
      - if it's an index I:
        - apply conjuncts in C that match I
        - apply rest of conjuncts to retrieved tuples
      - example
        - $c < 100 \text{ AND } a = 3 \text{ AND } b = 5$ 
          - can use a B+ tree index on  $c$  and check  $a = 3 \text{ AND } b = 5$  for each retrieved tuple
          - can use a hash index on  $a$  and  $b$  and check  $c < 100$  for each retrieved tuple

## Selection

- general selections - selections without disjunctions
  - evaluation options:
    2. use several indexes - when several conjuncts match indexes using a2 / a3
      - compute sets of rids of candidate tuples using indexes
      - intersect sets of rids, retrieve corresponding tuples
      - apply remaining conjuncts (if any)
      - example:  $c < 100 \text{ AND } a = 3 \text{ AND } b = 5$ 
        - use a B+ tree index on  $c$  to obtain rids of records that meet condition  $c < 100$  ( $R_1$ )
        - use a hash index on  $a$  to retrieve rids of records that meet condition  $a = 3$  ( $R_2$ )
        - compute  $R_1 \cap R_2 = R_{int}$
        - retrieve records with rids in  $R_{int}$  ( $R$ )
        - check  $b = 5$  for each record in  $R$

## Selection

- general selections
  - selections with disjunctions
- C - CNF condition with disjunctions, i.e., some conjunct  $J$  is a disjunction of terms
  - if some term  $T$  in  $J$  requires a file scan, testing  $J$  by itself requires a file scan
    - example:  $a < 100 \vee b = 5$ 
      - hash index on  $b$ , hash index on  $c$
  - => check both terms using a file scan (i.e., best access path: file scan)
- compare with the example below:
  - $(a < 100 \vee b = 5) \wedge c = 7$
  - hash index on  $b$ , hash index on  $c$
  - => use index on  $c$ , apply  $a < 100 \vee b = 5$  to each retrieved tuple (i.e., most selective access path: index)

## Selection

- general selections
  - selections with disjunctions
- C - CNF condition with disjunctions
  - every term  $T$  in a disjunction matches an index

=> retrieve tuples using indexes, compute union

  - example
    - $a < 100 \vee b = 5$
    - B+ tree indexes on  $a$  and  $b$
    - use index on  $a$  to retrieve records that meet condition  $a < 100$  ( $R_1$ )
    - use index on  $b$  to retrieve records that meet condition  $b = 5$  ( $R_2$ )
    - compute  $R_1 \cup R_2 = R$
    - if all matching indexes use a2 or a3 => take union of rids, retrieve corresponding tuples

# Projection

- $\Pi_{\text{SID, CID}}(\text{Exams})$

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- to implement projection:
  - eliminate:
    - unwanted columns
    - duplicates
- projection algorithms - techniques
  - sorting
  - hashing

## Projection Based on Sorting

- step 1
  - scan  $E \Rightarrow$  set of tuples containing only desired attributes ( $E'$ )
  - cost:
    - scan  $E$ :  $M$  I/Os
    - write temporary relation  $E'$ :  $T$  I/Os
      - $T$  depends on: number of columns and their sizes,  $T$  is  $O(M)$
- step 2
  - sort tuples in  $E'$
  - sort key – all columns
  - cost:  $O(T \log T)$  (also  $O(M \log M)$ )
- step 3
  - scan sorted  $E'$ , compare adjacent tuples, eliminate duplicates
  - cost:  $T$
- total cost:  $O(M \log M)$

## Projection Based on Sorting

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- scan Exams: 1000 I/O
- size of tuple in E': 10 bytes

=> cost of writing temporary relation E': 250 I/Os

- available buffer pages: 20
  - E' can be sorted in 2 passes
  - sorting cost:  $2 * 2 * 250 = 1000$  I/Os
- final scan of E' - cost: 250 I/Os

=> total cost: 2500 I/Os

\* E – record size = 40 bytes \*                      \* 1000 pages \*                      \* 100 records / page\*



## Projection Based on Sorting

- improvement
  - modify the sorting algorithm to do projection with duplicate elimination
    - modify pass 0 of external sort - eliminate unwanted columns
      - read in B pages from E
      - write out  $(T/M) * B$  internally sorted pages of E'
        - more aggressive approach: write out  $2*B$  internally sorted pages of E' (on average)
      - tuples in runs - smaller than input tuples
    - modify merging passes - eliminate duplicates
      - number of result tuples is smaller than number of input tuples


## Projection Based on Sorting

- improvement
  - pass 0
    - scan Exams: 1000 I/Os
    - write out 250 pages
    - 20 available buffer pages
      - 250 pages => 7 sorted runs about 40 pages long (except the last one)
  - pass 1
    - read in all runs – cost: 250 I/O
    - merge runs
- total cost : 1500 I/O


# Projection Based on Hashing

- phases
  - partitioning
  - duplicate elimination

original relation




disk



B pages in the buffer

partitions




disk

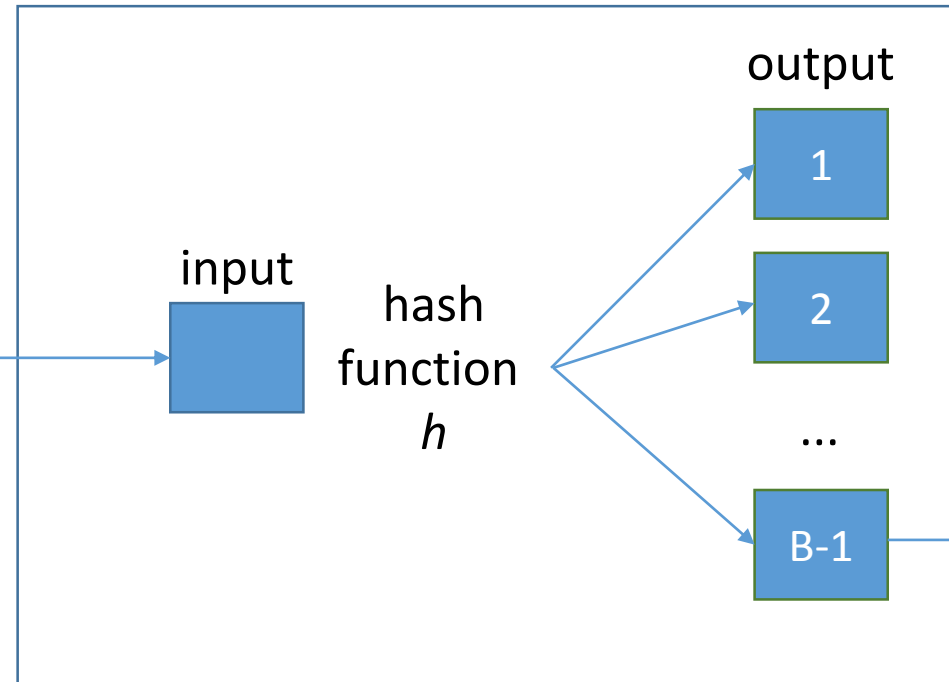
# Projection Based on Hashing

- partitioning phase
  - 1 input buffer page – read in the relation one page at a time
  - hash function  $h$  – distribute tuples uniformly to one of  $B-1$  partitions
  - $B-1$  output buffer pages – one output page / partition

original relation




disk



B pages in the buffer

partitions



disk

## Projection Based on Hashing

- partitioning phase
  - read the relation using the input buffer page
  - for each tuple  $t$ 
    - discard unwanted fields
    - apply hash function  $h$  to the combination of all remaining attributes
    - write  $t$  to the output buffer page that it is hashed to by  $h$

=> B-1 partitions

- partition
  - collection of tuples
    - common hash value
    - no unwanted fields
- 2 tuples in different partitions are guaranteed to be distinct

## Projection Based on Hashing

- duplicate elimination phase
  - process all partitions
    - read in partition P, one page at a time
      - build in-memory hash table for P with hash function  $h_2 (\neq h)$  on all fields
        - if a new tuple hashes to the same value as an existing tuple, compare them to check if they are distinct
        - eliminate duplicates as they are detected
      - write duplicate-free hash table to result file
      - clear in-memory hash table
  - partition overflow
    - apply hash-based projection technique recursively (subpartitions)

## Projection Based on Hashing

- cost
  - partitioning
    - read E: M I/Os
    - write E': T I/Os
  - duplicate elimination
    - read in partitions: T I/Os

=> total cost:  $M + 2 * T$  I/Os

- Exams:
  - $1000 + 2 * 250 = 1500$  I/Os

## Set Operations

- intersection, cross-product
  - special cases of join (i.e., join condition for intersection - equality on all fields, no join condition for cross-product)
- union, set-difference
  - similar
- union:  $R \cup S$ 
  - sorting
    - sort R and S on all attributes
    - scan the sorted relations in parallel; merge them, eliminating duplicates
    - refinement
      - produce sorted runs of R and S, merge runs in parallel



## Set Operations

- union:  $R \cup S$ 
  - hashing
    - partition  $R$  and  $S$  with the same hash function  $h$
    - for each  $S$ -partition
      - build in-memory hash table (using  $h_2$ ) for the  $S$ -partition
      - scan corresponding  $R$ -partition, add tuples to hash table, discard duplicates
      - write out hash table
      - clear hash table

## Aggregate Operations

- without grouping
  - scan relation
  - maintain *running information* about scanned tuples
    - COUNT - count of values retrieved
    - SUM - *total* of values retrieved
    - AVG -  $\langle total, count \rangle$  of values retrieved
    - MIN, MAX - smallest / largest value retrieved
- with grouping
  - sort relation on the grouping attributes
  - scan relation to compute aggregate operations for each group
  - improvement: combine sorting with aggregation computation
  - alternative approach based on hashing

## Aggregate Operations

- using existing indexes
  - index with a search key that includes all the attributes required by the query
    - index-only scan
  - attribute list in the GROUP BY clause is a prefix of the index search key (tree index)
    - get data entries (and records, if necessary) in the required order
    - i.e., avoid sorting

\* balanced merge sort\*

- table T,  $|T| = 3100$  records
- 1 run – at most 100 records
- runs distribution – 4 files
- initial runs – distribution
  - about half of the files
- merge runs, write runs to remaining files
- continue until a single run is produced
- notation
  - $x^y$ 
    - y runs with relative length = 1
  - run with relative length 1
    - produced with an internal sorting algorithm

\* balanced merge sort\*

F1	F2	F3	F4	obs
1 <sup>16</sup>	1 <sup>15</sup>	-	-	runs distribution
-	-	2 <sup>8</sup>	2 <sup>7</sup> 1 <sup>1</sup>	merging – alternate F3 and F4; copy one run
4 <sup>4</sup>	4 <sup>3</sup> 3 <sup>1</sup>	-	-	merging – alternate F1 and F2
-	-	8 <sup>2</sup>	8 <sup>1</sup> 7 <sup>1</sup>	merging – alternate F3 and F4
16 <sup>1</sup>	15 <sup>1</sup>	-	-	merging – alternate F1 and F2
-	-	31 <sup>1</sup>	-	merging, final result obtained in F3

## \* polyphase merge sort \*

- determine an initial configuration for the distribution of runs
  - one file  $F$  - empty
- merge runs from all files into  $F$  until a file  $F'$  is empty
- continue until a single run is produced

\* polyphase merge sort \*

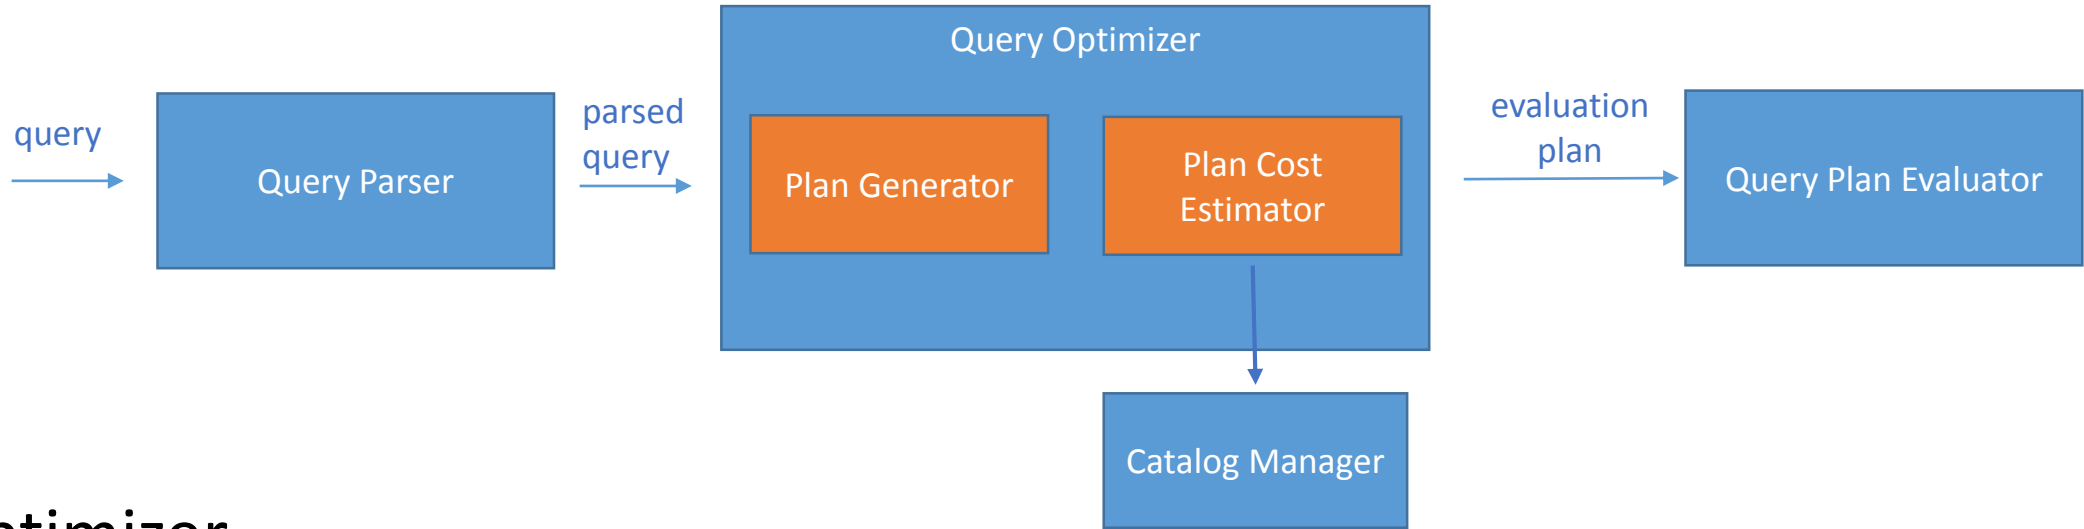
F1	F2	F3	F4	obs
-	1 <sup>7</sup>	1 <sup>11</sup>	1 <sup>13</sup>	runs distribution
3 <sup>7</sup>	-	1 <sup>4</sup>	1 <sup>6</sup>	merge in F1 until F2 is empty
3 <sup>3</sup>	5 <sup>4</sup>	-	1 <sup>2</sup>	merge in F2 until F3 is empty
3 <sup>1</sup>	5 <sup>2</sup>	9 <sup>2</sup>	-	merge in F3 until F4 is empty
-	5 <sup>1</sup>	9 <sup>1</sup>	17 <sup>1</sup>	merge in F4 until F1 is empty
31 <sup>1</sup>	-	-	-	merge, final result obtained in F1

\* polyphase merge sort \* - initial configuration

F1	F2	F3	F4	obs
1	-	-	-	the run in F1 is obtained from the other files
-	1	1	1	the run in F4 is obtained from the other files
1	2	2	-	runs in F3 are obtained from the other files
3	4	-	2	runs in F2 are obtained from the other files
7	-	4	6	runs in F1 are obtained from the other files
-	7	11	13	runs in F4 are obtained from the other files
13	20	24	-	runs in F3 are obtained from the other files
...	...	...	...	...
$c_n$	$b_n$	$a_n$	-	$a_n \geq b_n \geq c_n$
$c_n + a_n$	$b_n + a_n$	-	$a_n$	



# Query Optimization



- optimizer
  - objective
    - given a query  $Q$ , find a good evaluation plan for a  $Q$
  - generates alternative plans for  $Q$ , estimates their costs, and chooses the one with the least estimated cost
  - uses information from the system catalogs

- running example - schema
  - Students (SID: integer, SName: string, Age: integer)
  - Courses (CID: integer, CName: string, Description: string)
  - Exams (SID: integer, CID: integer, EDate: date, Grade: integer, FacultyMember: string)
- Students
  - every record has 50 bytes
  - there are 80 records / page
  - 500 pages
- Courses
  - every record has 40 bytes
  - there are 100 records / page
  - 1 page
- Exams
  - every record has 40 bytes
  - there are 100 records / page
  - 1000 pages


# Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      S.Age > 23
```

query

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

relational algebra expression

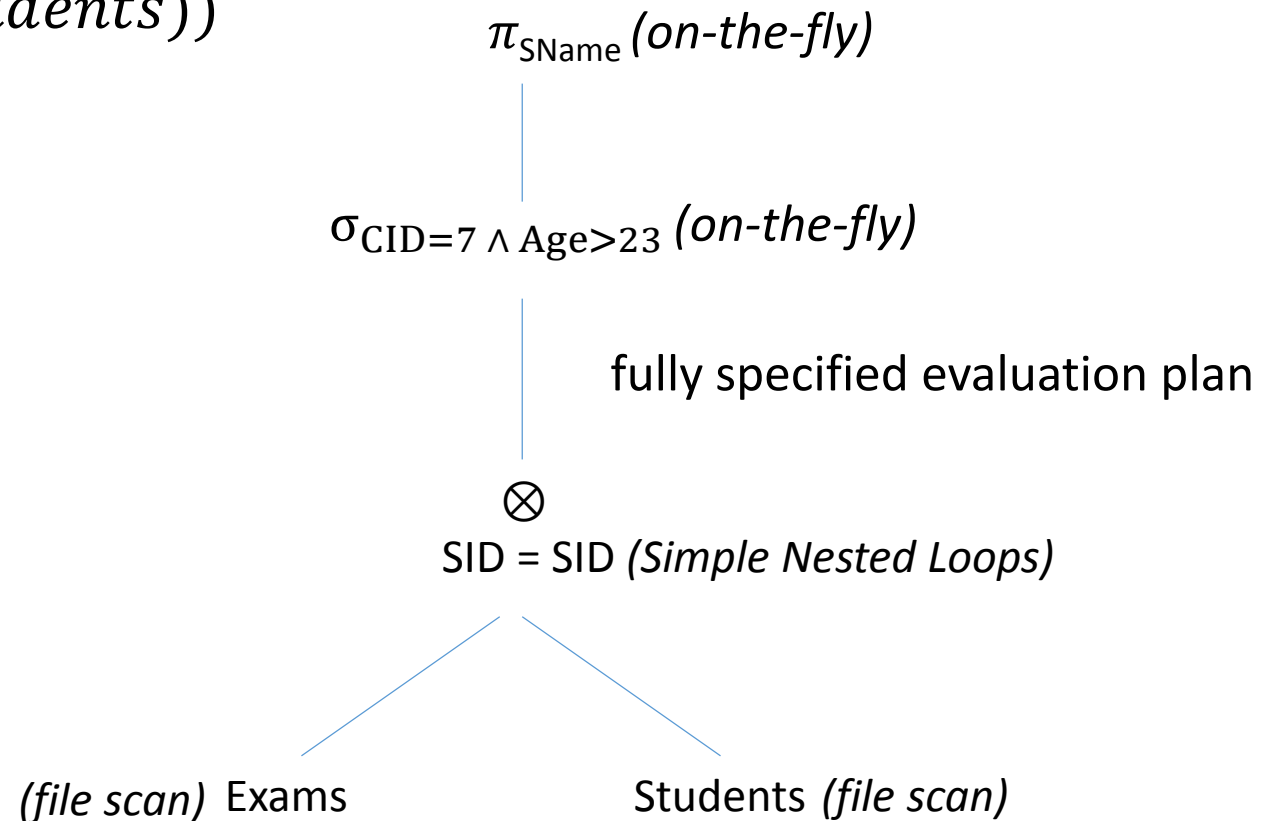


# Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      S.Age > 23
```

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

- query evaluation plan
  - extended relational algebra tree
  - node – annotations
    - relation
      - access method
    - relational operator
      - implementation method

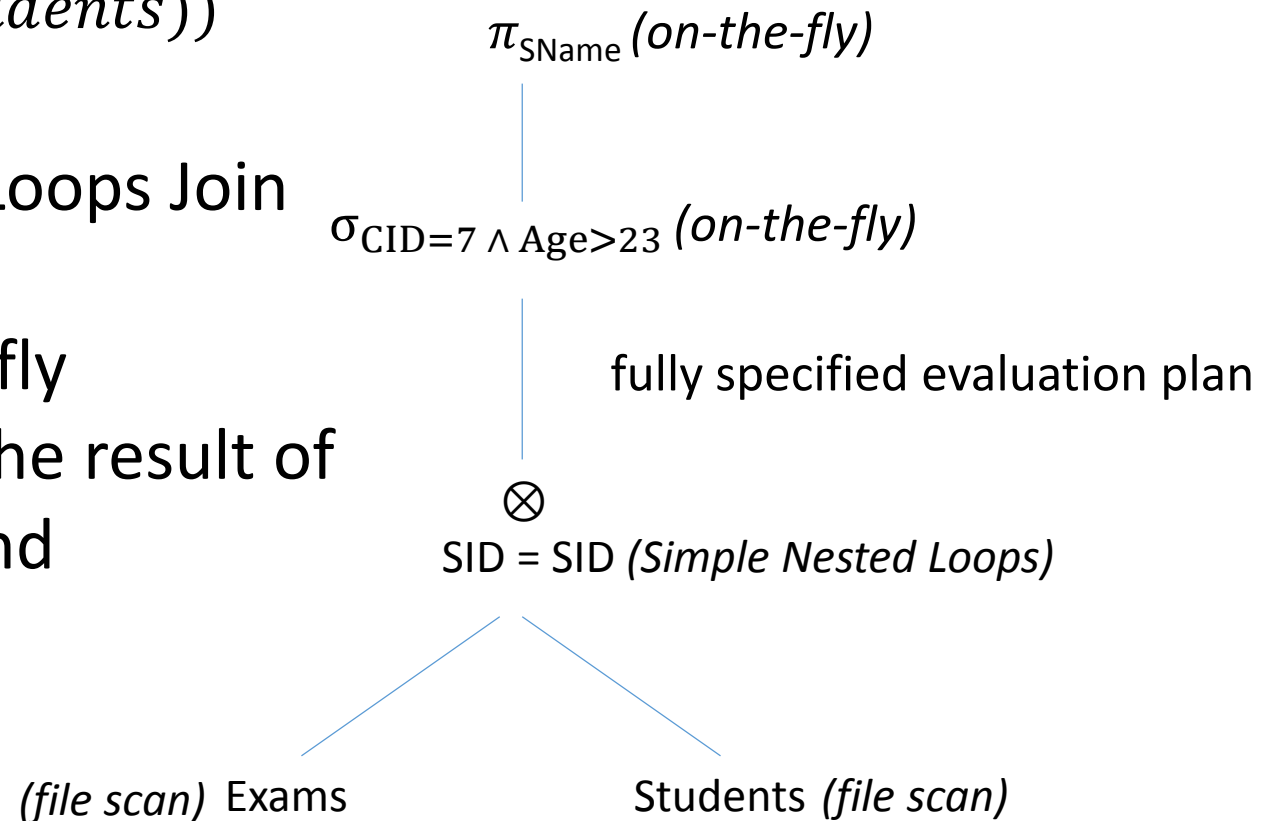


# Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      S.Age > 23
```

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

- e.g., page-oriented Simplified Nested Loops Join
- Exams – outer relation
- selection, projection applied on-the-fly to each tuple in the join result, i.e., the result of the join (before applying selection and projection) is not stored



# Pipelined Evaluation

SELECT \*

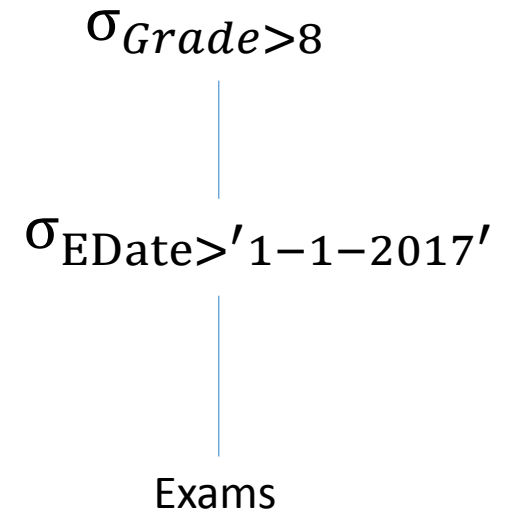
FROM Exams

WHERE EDate > '1-1-2017' AND Grade > 8

*T1*

*T2*

$\sigma_{Grade>8}(\sigma_{EDate>'1-1-2017'}(Exams))$



- index / matches *T1*
- v1 - *materialization*
  - evaluate *T1*
  - write out result tuples to temporary relation *R*, i.e., tuples are *materialized*
  - apply the 2<sup>nd</sup> selection to *R*
  - cost: read and write *R*

# Pipelined Evaluation

SELECT \*

FROM Exams

WHERE EDate > '1-1-2017' AND Grade > 8

*T1*

*T2*

$\sigma_{Grade > 8}$

$\sigma_{EDate > '1-1-2017'}$

Exams

- *v2 – pipelined evaluation*
  - apply the 2<sup>nd</sup> selection to each tuple in the result of the 1<sup>st</sup> selection as it is produced
  - i.e., 2<sup>nd</sup> selection operator is applied *on-the-fly*
  - saves the cost of writing out / reading in the temporary relation *R*
- iterator interface
  - combining code for individual operators into an executable plan (functions *open, get\_next, close*)
  - supports pipelining

## Query Blocks – Units of Optimization

- optimize SQL query Q
  - parse Q => collection of query *blocks*
  - optimizer:
    - optimize one block at a time
- query *block* - SQL query:
  - without nesting
  - with exactly: one SELECT clause, one FROM clause
  - with at most: one WHERE clause, one GROUP BY clause, one HAVING clause
    - WHERE condition - CNF



# Query Blocks – Units of Optimization

- query Q:

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = (SELECT MAX(S2.Age)
               FROM Students S2)
GROUP BY S.SID
HAVING COUNT(*) > 2
```

outer block

nested block

- decompose query into a collection of blocks without nesting

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

# Query Blocks – Units of Optimization

## \* block optimization

- express query block as a relational algebra expression

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

$$\pi_{S.SID, MIN(E.EDate)}($$
$$HAVING_{COUNT(*) > 2}($$
$$GROUP BY_{S.SID}($$
$$\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value\_from\_nested\_block}($$
$$Students \times Exams \times Courses)))))$$

- GROUP BY, HAVING – operators in the extended algebra used for plans
- argument list of projection can include aggregate operations

## Query Blocks – Units of Optimization

- query Q treated as a  $\sigma \pi \times$  algebra expression
- the remaining operations in Q are performed on the result of the  $\sigma \pi \times$  expression

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

$$\pi_{S.SID, E.EDate}(\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value\_from\_nested\_block}(Students \times Exams \times Courses))$$

- attributes in GROUP BY, HAVING are added to the argument list of projection
- aggregate expressions in the argument list of projection are replaced by their argument attributes

## Query Blocks – Units of Optimization

### \* block optimization

- find best plan  $P$  for the  $\sigma \pi \times$  expression
- evaluate  $P \Rightarrow$  result set  $RS$
- sort/hash  $RS \Rightarrow$  groups
- apply HAVING to eliminate some groups
- compute aggregate expressions in SELECT for each remaining group

$\pi_{S.SID, MIN(E.EDate)}(\$   
 $HAVING_{COUNT(*) > 2}(\$   
 $GROUP BY_{S.SID}(\$   
 $\pi_{S.SID, E.EDate}(\$   
 $\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value\_from\_nested\_block}(\$   
 $Students \times Exams \times Courses ))))$

# References

- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2<sup>nd</sup> Edition), McGraw-Hill, 2000
- [Da03] DATE, C.J., An Introduction to Database Systems (8<sup>th</sup> Edition), Addison-Wesley, 2003
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007, <http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si10] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, McGraw-Hill, 2010, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems, <http://infolab.stanford.edu/~ullman/fcdb.html>